PHASE 1: FULL REPOSITORY INTEGRATION & CONFLICT AUDIT

> **AUDIT DIRECTIVE COMPLIANCE NOTICE**
> In strict accordance with the Phase 1 instructions, **zero code files, database schemas, API routes, or configuration files were modified, moved, renamed, deleted, or refactored**. This is an exhaustive, factual, read-only audit of the coexistence state between the **Local AI Foundation Architecture** (core/, connectors/, orchestration/, workflows/, apps/, alembic/, configs/) and the **Standalone RAG Subsystem** (rag/).

---

1. Repository Inventory & Categorization Map

The repository at /home/vanilla/Desktop/hackathon/sih/local-ai contains two major subsystem architectures that coexist on the active branch anish.

```
local-ai/
├── adapters/                [Foundation: Hardware/Engine Adapters]
│   └── llama_cpp/           [Uninitialized Git submodule for llama.cpp]
├── alembic/                 [Foundation: Database Migrations]
│   ├── versions/
│   │   └── 001_initial_orchestration_schema.py
│   └── env.py
├── apps/                    [Foundation: Delivery / Application Layer]
│   ├── api/                 [FastAPI Delivery Server]
│   │   ├── routers/         [artifacts.py, direct.py, files.py, goals.py, telemetry.py]
│   │   ├── schemas/         [Pydantic request/response schemas]
│   │   ├── app.py           [FastAPI application factory, CORS, RFC 7807 problem handlers]
│   │   ├── dependencies.py  [FastAPI DI injection for AppContext]
│   │   └── events.py        [EventBus for SSE streaming]
│   └── context.py           [AppContext: Container wiring core + orchestration engines]
├── artifacts/               [Foundation: Generated artifact outputs (.xlsx, .docx, .pdf)]
├── configs/                 [Foundation: Configuration files]
│   ├── llama_models.ini     [Model presets for llama-server router mode]
│   ├── models/              [YAML model profiles]
│   └── settings.toml        [Foundation TOML configuration]
├── connectors/              [Foundation: Architectural Seams / Capability Boundaries]
│   ├── base.py              [Connector interfaces]
│   ├── factory.py           [Connector factory]
```

```
|   └── inference.py          [FoundationInferenceConnector]
├── core/                     [Foundation: Low-Level Core Platform]
|   ├── common/               [errors.py, parsing.py, formatting.py]
|   ├── config/               [Pydantic settings loaders]
|   ├── inference/            [Inference contracts, TokenUsage, StreamingChunk]
|   ├── models/               [Model registry, ModelProfile]
|   └── providers/            [ProviderManager, LlamaCppProvider]
├── docs/                     [Foundation & System Documentation]
├── models/                   [Foundation: Model storage directory]
|   └── gguf/                 [Target directory for .gguf model weights]
├── orchestration/            [Foundation: Agentic & Task Orchestration Engine]
|   ├── capabilities/         [Builtin capabilities: agent, artifact, code, doc, vision, etc.]
|   ├── decision/             [Replanning engine, task failure decisioning]
|   ├── persistence/          [PostgreSQL / SQLite persistence, engine.py, repository.py]
|   ├── planning/             [Planners, DAG generators, replanning]
|   ├── routing/              [Staged router, Aurelio semantic router, LLM classifier]
|   ├── runners/              [InProcessTaskRunner, TemporalRunner]
|   ├── temporal/             [Temporal workflow definitions and activities]
|   └── validation/           [Plan and DAG acyclicity validators]
├── rag/                      [RAG Subsystem: Standalone Modular Ingestion & Retrieval Pipeline]
|   ├── chunking/             [Structural recursive chunker, Chunk dataclass]
|   ├── cli/                  [console.py developer test harness & verification tool]
|   ├── embedding/            [NomicEmbedder, ChunkEmbeddingService, EmbeddingResult]
|   ├── indexing/             [pgvector persistence, VectorIndexer]
|   ├── ingestion/            [DoclingDocumentIngester, IngestedDocument, IngestionResult]
|   ├── metadata/             [ChunkMetadata, ProvenanceTracker]
|   ├── normalization/        [Text normalizer, table/code/markdown cleanup]
|   ├── rerank/               [CrossEncoderReranker, CandidateChunk, RerankResult]
|   ├── retrieval/            [VectorRetriever, VectorSearchResult, pgvector similarity search]
|   ├── storage/              [database.py DatabaseManager, models.py DocumentModel & ChunkModel]
|   ├── tests/                [154 unit & integration tests covering RAG pipeline]
|   └── offline.py            [Zero-internet / air-gapped environment enforcement]
├── scripts/                  [Foundation: Shell build & launch scripts]
|   ├── build_llama_cpp.sh    [CUDA Ninja build script for llama.cpp]
|   ├── run_smoke_test.sh     [Smoke test script for llama-cli]
|   └── start_llama_server.sh [Launch script for llama-server with llama_models.ini]
├── tests/                    [Foundation: Unit and integration test suite]
├── workflows/                [Foundation: Specialized Domain Workflows]
|   ├── code_repair/          [Automated workspace code repair workflow]
|   └── text_analysis.py      [Structured text analysis two-pass workflow]
├── .env                      [RAG runtime environment file]
├── alembic.ini               [Alembic configuration]
└── pyproject.toml            [Project build & dependency declaration]
```

---

2. Git & History Analysis

- **Current Active Branch**: anish
- **Head Commit**: d7ed39f (Merge branch 'master' into anish)
- **Integration Merge Details**:
  - master branch (carrying the complete Foundation Architecture built by the second developer) was merged directly into anish (carrying the completed RAG pipeline).
  - The merge touched **182 files** with **27,428 insertions** and **0 deletions**.
  - **No git merge conflict occurred** because the Foundation was created in new namespaces (core/, orchestration/, connectors/, apps/, workflows/, alembic/, configs/), leaving rag/ completely isolated in its own directory tree.
- **Divergence Risk**:
  - While there were zero source-file git merge conflicts, **semantic divergence** exists in configuration, database targets, dependency declarations, and document ingestion models.

---

3. Complete 16-Endpoint Table + API Analysis

The delivery application in apps/api/ defines exactly 16 HTTP endpoints across 5 router modules, mounted under prefix /api/v1.

Endpoint Inventory

#	HTTP Method	Route Path	Router File	Handler Function	Request Schema / Params	Response Schema	Functionality
1	POST	/api/v1/files/upload	apps/api/routers/files.py	upload_file	UploadFile (multipart/form-data)	FileMetadataResponse	Uploads file to local disk store, validates MIME, returns SHA256 & file ID
2	GET	/api/v1/files/{file_id}	apps/api/routers/files.py	get_file	file_id: str	FileResponse	Downloads raw uploaded file content by ID
3	DELETE	/api/v1/files/{file_id}	apps/api/routers/	delete_file	file_id: str	204 No Content	Removes file from

#	HTTP Method	Route Path	Router File	Handler Function	Request Schema / Params	Response Schema	Functionality
			files.py				disk store and tracking registry
4	POST	/api/v1/direct/chat	apps/api/routers/direct.py	direct_chat	DirectChatRequest	DirectChatResponse	Direct inference call through InferenceConnector without orchestrato planning
5	POST	/api/v1/direct/documents	apps/api/routers/direct.py	direct_documents	DirectDocumentRequest	DirectDocumentResponse	Executes document.understand capability directly via DoclingDocumentParser
6	POST	/api/v1/direct/vision	apps/api/routers/direct.py	direct_vision	DirectVisionRequest	DirectVisionResponse	Executes vision.analyze capability directly on image payloads
7	POST	/api/v1/direct/artifacts	apps/api/routers/direct.py	direct_artifacts	DirectArtifactRequest	DirectArtifactResponse	Executes artifact.generate capability to create Excel, Word, or PDF deliverables
8	POST	/api/v1/direct/workspaces	apps/api/routers/direct.py	direct_workspaces	DirectWorkspaceRequest	DirectWorkspaceResponse	Executes sandboxed code execution in Docker workspace container
9	POST	/api/v1/goals	apps/api/routers/goals.py	submit_goal	GoalSubmitRequest	GoalResponse	Submits natural language goal, generates DAG plan starts execution

#	HTTP Method	Route Path	Router File	Handler Function	Request Schema / Params	Response Schema	Functionality
10	GET	/api/v1/goals/{goal_id}	apps/api/routers/goals.py	get_goal	goal_id: str	GoalResponse	Fetches goal aggregate status, active plan, tasks, and deliverables
11	GET	/api/v1/goals/{goal_id}/stream	apps/api/routers/goals.py	stream_goal	goal_id: str	text/event-stream (SSE)	Real-time Server-Sent Events stream for task lifecycle & log progression
12	POST	/api/v1/goals/{goal_id}/cancel	apps/api/routers/goals.py	cancel_goal	goal_id: str	GoalResponse	Signals cancellation to orchestrator, aborting running and pending tasks
13	GET	/api/v1/goals/{goal_id}/plans	apps/api/routers/goals.py	get_goal_plans	goal_id: str	List[PlanResponse]	Retrieves historical and revised execution plans for a given goal
14	GET	/api/v1/artifacts/{artifact_id}	apps/api/routers/artifacts.py	get_artifact	artifact_id: str	ArtifactMetadataResponse	Retrieves metadata, checksum, mime type, and location of generated artifact
15	GET	/api/v1/artifacts/{artifact_id}/download	apps/api/routers/artifacts.py	download_artifact	artifact_id: str	FileResponse	Downloads binary deliverable file (.xlsx, .docx, .pdf, .zip)
16	GET	/api/v1/telemetry/stats	apps/api/routers/telemetry.py	get_telemetry_stats	None	TelemetryStatsResponse	In-memory token counts, latency percentiles, and provider error stats

API Gap & Collision Analysis

1. **Zero Endpoint Collisions**: There is currently **not a single endpoint** in `apps/api/` representing RAG functionality (no chunking, embedding, vector search, or semantic retrieval routes).
2. **Missing RAG API Surfaces**:
    - Ingestion Route: No endpoint exists to ingest and index an uploaded document into the RAG vector store.
    - Retrieval Route: No endpoint exists to retrieve relevant semantic chunks or reranked candidates for a query.
    - Direct RAG Query Route: No endpoint exists to perform Adaptive RAG or grounded QA.
3. **Seam Opportunity**: The Foundation has `/api/v1/direct/documents` which currently parses documents for raw text/table extraction using Docling without vector persistence. RAG can cleanly add `/api/v1/rag/ingest` and `/api/v1/rag/retrieve` or integrate behind `/api/v1/direct/retrieve`.

---

4. RAG ↔ Foundation Compatibility & Adapter Architecture

Data Models Comparison

Dimension	Foundation Concept	RAG Pipeline Concept	Divergence & Reconciliation
Document Ingestion	`DoclingDocumentParser` (`NormalizedDocument`, `ExtractedTable`, `ExtractedFigure`)	`DoclingDocumentIngester` (`IngestedDocument`, `DocumentModel`)	Both use Docling 2.126. Foundation extracts high-level structure; RAG extracts structured markdown with headings and layout hierarchies tailored for recursive chunking.
Chunk Model	`DataReference` (generic pointer to content/file)	`Chunk` (`ChunkMetadata`, `ChunkModel`)	Foundation treats data inputs as reference URIs. RAG treats chunks as discrete semantic units with 768-dim embeddings, byte provenance, page numbers, and section breadcrumbs.
Output Artifacts	`ArtifactReference` / `Artifact` (disk path, sha256, mime type)	Vector store records in `rag_documents` and `rag_chunks`	RAG produces indexed embeddings in PostgreSQL; Foundation produces reports (.xlsx, .docx, .pdf). RAG

Dimension	Foundation Concept	RAG Pipeline Concept	Divergence & Reconciliation
			retrieval can feed directly into Foundation artifact generators.
System Context	AppContext (wires FoundationCore, Orchestrator, InProcessTaskRunner)	DatabaseManager, VectorIndexer, VectorRetriever, CrossEncoderReranker	AppContext can hold a single RAGService or RetrievalConnector instance, initializing RAG components lazily.

Architectural Seam Analysis

The Foundation was designed around the **Connector Pattern**:

- connectors/base.py: Defines abstract connector boundaries.
- connectors/inference.py: Implements FoundationInferenceConnector to decouple workflows from raw LLM providers.
- orchestration/capabilities/builtin/: Registers built-in task capabilities (agent, artifact, code, code_repair, document, inference, vision, workflow).

Clear Architectural Seam: RAG fits naturally into the Foundation at two well-defined seams **without changing any internal RAG code**:

1. **At the Connector Layer**: A new RetrievalConnector interface in connectors/retrieval.py with an implementation LocalRAGRetrievalConnector wrapping rag.retrieval.VectorRetriever and rag.rerank.CrossEncoderReranker.
2. **At the Capability Layer**: A new built-in capability retrieval.search and rag.ingest registered in orchestration/capabilities/builtin/retrieval/ allowing the Planner and Orchestrator to schedule semantic search tasks in execution plans.

---

5. Inference Conflict Audit

A critical question was whether RAG and the Foundation compete or collide over generative inference:

Dimension	Local AI Foundation	RAG Subsystem (rag/)	Conflict Status
Inference Framework	FoundationCore + ProviderManager + LlamaCppProvider	None (No LLM generation implemented yet)	**NO CONFLICT**
Inference Server	External llama-server on http://127.0.0.1:8080 (Qwen3.5-9B, Qwen3.5-0.8B GGUF)	None	**NO CONFLICT**
Embedding Engine	None	Hugging Face Transformers (nomic-ai/nomic-embed-text-v1.5, PyTorch CPU/GPU, 768-dim)	**NO CONFLICT**
Reranking Engine	None	Sentence-Transformers (cross-encoder/ms-marco-MiniLM-L-6-v2, PyTorch CPU)	**NO CONFLICT**
Tokenizers	Core inference tokenizers for llama context budgeting	HF AutoTokenizer for Nomic & CrossEncoder token management	**NO CONFLICT**

Findings

- **Zero LLM Inference Duplication**: rag/ does not implement any LLM inference, prompt generation, or answer synthesis.
- **Hardware Resource Coexistence**:
  - llama-server runs on GPU (VRAM ~6-8GB for Q4_K_M).
  - Nomic Embeddings (nomic-embed-text-v1.5) runs locally on PyTorch (~560MB memory).
  - CrossEncoder (ms-marco-MiniLM-L-6-v2) runs locally on CPU/PyTorch (~130MB memory).
  - Both fit comfortably on a single machine with >= 16GB RAM and an 8GB-12GB GPU.

---

6. Database / PostgreSQL & Schema Coexistence Audit

Divergence Analysis

Parameter	Foundation Persistence	RAG Subsystem Storage	Conflict / Coexistence
Database Target	local_ai (configured in configs/settings.toml & alembic.ini)	local_ai_rag (default in rag/storage/database.py) OR ragdb (in .env)	**Configuration Divergence**: Three different database names currently configured across settings!
Database Driver	postgresql+psycopg:// (Psycopg 3)	postgresql+psycopg:// (Psycopg 3)	**Identical driver** (100% compatible)
Connection Pooling	orchestration.persistence.engine.create_db_engine()	rag.storage.database.DatabaseManager()	Two independent engine/pool instances
Extensions Required	None (Standard PostgreSQL JSONB & relational)	CREATE EXTENSION IF NOT EXISTS vector; (pgvector)	RAG requires pgvector installed on the PostgreSQL instance

Table Name Collision Analysis

Foundation Schema (orchestration_*)	RAG Subsystem Schema (rag_*)	Collision?
orchestration_goals	rag_documents	**NONE**
orchestration_plans	rag_chunks (with 768-dim vector column)	**NONE**

Foundation Schema (orchestration_*)	RAG Subsystem Schema (rag_*)	Collision?
orchestration_plan_revisions	—	**NONE**
orchestration_tasks	—	**NONE**
orchestration_dependencies	—	**NONE**
orchestration_attempts	—	**NONE**

Schema Migration Alignment

- Foundation uses **Alembic** (alembic/versions/001_initial_orchestration_schema.py).
- RAG currently uses **declarative initialization** (rag.storage.database.DatabaseManager.init_db() which executes Base.metadata.create_all()).
- Because all Foundation tables are prefixed with orchestration_ and all RAG tables are prefixed with rag_, they can **safely reside in the exact same database** (e.g. local_ai) without collision. A second Alembic migration (002_rag_vector_schema.py) can bring RAG tables under formal Alembic revision management.

---

7. Configuration Inventory & Ownership Table

Configuration File	Owner Subsystem	Purpose & Managed Keys	Conflicts & Overlaps
configs/settings.toml	Foundation	[foundation], [providers.llama_cpp], [database] url="…/local_ai", [document], [artifact], [workspace]	Points to database local_ai. Does not contain RAG settings (embedding models, pgvector dimension, reranker model).
configs/llama_models.ini	Foundation	Presets for llama-server: context size, batch size, cache types, model paths for qwen3.5-9b and qwen3.5-0.8b	Pure inference configuration. No overlap with RAG.
.env	RAG	RAG_DATABASE_URL=postgresql+psycopg://raguser:ragpassword@localhost:5432/ragdb	**Direct Conflict with settings.toml**: Different user (raguser vs postgres), different password, and different db

Configuration File	Owner Subsystem	Purpose & Managed Keys	Conflicts & Overlaps
			name (ragdb vs local_ai).
alembic.ini	Foundation	sqlalchemy.url = postgresql+psycopg://postgres:postgres@localhost:5432/local_ai	Matches settings.toml. Conflicts with .env.
PROJECT_CONTEXT.md	Foundation	Architecture specification for MRPL sovereign on-premise workbench	Foundation developer design documentation.
mrpl.md	Foundation	Comprehensive MRPL design document	Foundation developer design documentation.

---

8. Infrastructure & Service Dependency Graph

Infrastructure Service Details

1. **PostgreSQL**: Single Postgres container or service with pgvector enabled. Needs user/db alignment.
2. **llama-server**: Required for Foundation direct chat, agent capabilities, and plan generation. Requires adapters/llama_cpp compilation and GGUF model files.
3. **Local Embedding / Reranking**: In-process PyTorch models loaded directly from disk cache (~/.cache/huggingface/hub/). No separate server process required.
4. **FastAPI**: Runs via uvicorn apps.api.app:app --host 0.0.0.0 --port 8000.
5. **Temporal**: Optional; the orchestrator has an InProcessTaskRunner fallback that executes tasks without a live Temporal cluster.

---

9. Offline / Air-Gapped Compatibility Analysis

The MRPL project strictly requires **ZERO runtime internet access**:

RAG Subsystem Status: FULLY AIR-GAPPED & HARDENED

- Module `rag/offline.py` configures:
    - `HF_HUB_OFFLINE=1`
    - `TRANSFORMERS_OFFLINE=1`
    - `HF_DATASETS_OFFLINE=1`
    - Socket connect blocking via monkeypatching in offline mode.
    - Docling: `enable_remote_services=False` explicitly passed to `DocumentConverter`.
    - Hugging Face cache audited:
        - `models--nomic-ai--nomic-embed-text-v1.5` exists in `~/.cache/huggingface/hub/`
        - `models--nomic-ai--nomic-bert-2048` exists in `~/.cache/huggingface/hub/`
        - `models--cross-encoder--ms-marco-MiniLM-L-6-v2` exists in `~/.cache/huggingface/hub/`
        - `models--docling-project--docling-layout-heron` exists in `~/.cache/huggingface/hub/`
        - `models--docling-project--docling-models` exists in `~/.cache/huggingface/hub/`
- Verified: All 135 unit & offline tests pass with zero network access.

Foundation Subsystem Status: PARTIALLY AIR-GAPPED (NEEDS HARDENING)

- **Vulnerability Identified**: In `orchestration/capabilities/builtin/document/docling_parser.py`, the `DocumentConverter` is instantiated without explicitly passing `enable_remote_services=False` or setting `offline.py` environment variables.
- If invoked in an environment with partial network connectivity, Docling's default OCR pipeline might attempt to verify remote Hugging Face model metadata unless protected.
- **Remediation**: The Foundation should reuse `rag/offline.py` or adopt its environment configuration at startup.

---

10. Dependency & pyproject.toml Audit

A substantial discrepancy was discovered between `pyproject.toml` and the virtual environment `.venv`:

Divergence Inventory

Package	Declared in pyproject.toml?	Installed in .venv?	Subsystem Dependent	Impact
fastapi	Yes (>=0.115.0)	**NO**	Foundation API	Foundation API cannot

Package	Declared in pyproject.toml?	Installed in .venv?	Subsystem Dependent	Impact
			(apps/api/)	launch in current .venv
uvicorn	Yes (>=0.30.0)	**NO**	Foundation API	Foundation ASGI server cannot launch
docker	Yes (>=7.0.0)	**NO**	Foundation Workspace Execution	Workspace container isolation fails
temporalio	Yes (>=1.32.0)	**NO**	Foundation Temporal Runner	Temporal workflows cannot run
semantic-router	Yes (>=0.1.16)	**NO**	Foundation Routing	Aurelio semantic router fails
pydantic-ai-slim	Yes (~=2.40.0)	**NO**	Foundation Agent Capability	PydanticAI model adapter fails
pgvector	**NO**	Yes (0.5.0)	RAG Storage & Indexing	Undeclared dependency in pyproject.toml
torch	**NO**	Yes (2.14.0)	RAG Embedding & Reranking	Undeclared dependency in pyproject.toml
transformers	**NO**	Yes (5.16.1)	RAG Embedding (Nomic)	Undeclared dependency in pyproject.toml
sentence-transformers	**NO**	Yes (6.0.1)	RAG Reranking (CrossEncoder)	Undeclared dependency in pyproject.toml
python-dotenv	**NO**	Yes (1.2.3)	RAG Storage Database	Undeclared dependency in pyproject.toml
rich	**NO**	Yes (15.0.0)	RAG CLI Console	Undeclared dependency in pyproject.toml
docling	Yes (>=2.120.0)	Yes (2.126.0)	Both (Shared)	Compatible

Package	Declared in pyproject.toml?	Installed in .venv?	Subsystem Dependent	Impact
sqlalchemy	Yes (>=2.0.0)	Yes (2.0.52)	Both (Shared)	Compatible
psycopg	Yes (>=3.1.0)	Yes (3.3.5)	Both (Shared)	Compatible

Pytest Configuration Gap

In pyproject.toml:

```
[tool.pytest.ini_options]
testpaths = ["tests"]
```

Because testpaths is hardcoded to ["tests"], any standard pytest run completely ignores all 154 tests under rag/tests/.

---

11. Layering & Dependency-Direction Audit

Foundation Layering Model

The Foundation is organized into a strict dependency hierarchy: $$\text{core} \longrightarrow \text{connectors} \longrightarrow \text{orchestration} \longrightarrow \text{workflows} \longrightarrow \text{apps}$$

1. core: Pure infrastructure, zero dependencies on outer layers.
2. connectors: Abstract protocol seams decoupling execution from implementations.
3. orchestration: Domain models (Goals, Plans, Tasks, DAGs) and execution engines.
4. workflows: Specialized multi-step pipelines (e.g. text_analysis.py, code_repair).
5. apps: FastAPI HTTP delivery layer and CLI entry points.

RAG Subsystem Layering Model

rag/ is a self-contained, modular vertical: $$\text{ingestion} \longrightarrow \text{normalization} \longrightarrow \text{chunking} \longrightarrow \text{metadata} \longrightarrow \text{embedding} \longrightarrow \text{indexing} \longrightarrow \text{retrieval} \longrightarrow \text{reranking}$$ Storage models (rag/storage/) and offline helpers (rag/offline.py) serve as internal utilities.

Dependency Direction Rules for Integration

1. **Rule 1**: rag/ MUST NOT import anything from core, orchestration, workflows, or apps. RAG remains an independent, reusable library.
2. **Rule 2**: connectors/ should introduce a RetrievalConnector protocol.
3. **Rule 3**: orchestration/capabilities/builtin/ can consume RetrievalConnector or import rag/ directly as an internal engine.
4. **Rule 4**: apps/api/ accesses RAG capabilities either via AppContext.orchestrator or through a dedicated retrieval service.

---

12. Entry Point & Startup Audit

Entry Point	Location	Execution Mechanism	Current Status / Issues
RAG Developer CLI	rag/cli/console.py	.venv/bin/python -m rag.cli.console	**READY & WORKING**. Interactive Rich terminal harness for PDF ingestion, chunk inspection, and semantic search.
Foundation API Server	apps/api/app.py	uvicorn apps.api.app:app --port 8000	**BLOCKED**. fastapi and uvicorn are not installed in the .venv.
llama-server Daemon	scripts/start_llama_server.sh	./scripts/start_llama_server.sh	**BLOCKED**. llama-server binary not compiled; GGUF models missing in models/gguf/.
llama.cpp Build Script	scripts/build_llama_cpp.sh	./scripts/build_llama_cpp.sh	**UNINITIALIZED**. adapters/llama_cpp submodule is empty; requires CUDA compiler /usr/local/cuda-12.8/bin/nvcc.
Inference Smoke Test	scripts/run_smoke_test.sh	./scripts/run_smoke_test.sh	**BLOCKED**. Requires llama-cli and Qwen3.5-9B-Q4_K_M.gguf.

---

13. Full-System Run Analysis

To execute the complete integrated system end-to-end, the following services must be prepared and launched:

Step 1: PostgreSQL with pgvector

```
# Must create single unified database 'local_ai' with vector extension
podman run -d --name local-ai-postgres \
  -e POSTGRES_USER=postgres \
  -e POSTGRES_PASSWORD=postgres \
  -e POSTGRES_DB=local_ai \
  -p 5432:5432 \
  pgvector/pgvector:pg16
```

(Status: PostgreSQL is present on the host system, but needs database alignment between `local_ai` and `local_ai_rag`).

Step 2: Compile and Launch llama-server **[MISSING PREREQUISITES]**

```
# 1. Initialize submodule
git submodule update --init --recursive adapters/llama_cpp  # [MISSING: submodule uninitialized]

# 2. Compile with CUDA
./scripts/build_llama_cpp.sh                          # [MISSING: llama-server binary]

# 3. Download GGUF models
# Place Qwen3.5-9B-Q4_K_M.gguf in models/gguf/            # [MISSING: GGUF model files]

# 4. Launch router
./scripts/start_llama_server.sh
```

Step 3: Run Database Migrations

```
alembic upgrade head
```

Step 4: Launch FastAPI Delivery Application **[MISSING PREREQUISITES]**

```
pip install fastapi uvicorn pydantic-ai-slim semantic-router
uvicorn apps.api.app:app --host 0.0.0.0 --port 8000        # [MISSING: dependencies in .venv]
```

Step 5: Run RAG Verification Harness

```
.venv/bin/python -m rag.cli.console                    # [READY]
```

---

14. Merge Classification (A-F)

Every overlapping or interfacing component in the repository is classified below:

Classification	Meaning
A: Keep As-Is	Do not touch; completely functional and self-contained
B: Adapter Boundary	Connect through an explicit interface/connector seam without altering internal code
C: Unify / Consolidate	Harmonize duplicated configuration or disparate database settings
D: Replace / Deprecate	Subsumed by a superior implementation
E: Needs Refactor	Code requires targeted adjustments to resolve a technical gap
F: Deferred	Optional enhancement postponed to post-integration phase

Detailed Classifications

1. **RAG Core Pipeline (`rag/chunking/`, `rag/embedding/`, `rag/indexing/`, `rag/metadata/`, `rag/normalization/`, `rag/rerank/`, `rag/retrieval/`, `rag/offline.py`)**:
    - **Classification: A (Keep As-Is)**
    - Rationale: High test coverage (154 tests), zero network leaks, self-contained, model-agnostic.
2. **RAG Developer CLI (`rag/cli/console.py`)**:
    - **Classification: A (Keep As-Is)**
    - Rationale: Fully functional standalone testing harness for PDF ingestion and vector retrieval.
3. **Database Target & Connection Management**:
    - **Classification: C (Unify / Consolidate)**
    - Rationale: Unify `local_ai_rag` / `ragdb` / `local_ai` into `local_ai`. Ensure pgvector extension and both table sets (`orchestration_*` and `rag_*`) reside in `local_ai`.
4. **Foundation Inference Stack (`core/inference/`, `core/providers/`, `llama-server`)**:
    - **Classification: A (Keep As-Is)**
    - Rationale: Production-grade LLM inference architecture.
5. **Document Ingestion (`rag/ingestion/docling.py` vs

orchestration/capabilities/builtin/document/docling_parser.py):

- **Classification: B (Adapter Boundary)**
- Rationale: Keep both specialized implementations; adapt rag/ingestion/ behind a retrieval indexing capability, while keeping docling_parser.py for direct document-to-JSON understanding. Add enable_remote_services=False to docling_parser.py.

6. **Retrieval Integration Seam**:
   - **Classification: B (Adapter Boundary)**
   - Rationale: Create connectors/retrieval.py and orchestration/capabilities/builtin/retrieval/ to expose RAG to the Planner and Agent.
7. **pyproject.toml and .venv Dependencies**:
   - **Classification: C (Unify / Consolidate)**
   - Rationale: Add missing RAG dependencies (pgvector, sentence-transformers, torch, transformers, rich, python-dotenv) to pyproject.toml. Add rag/tests to testpaths.

---

15. What Must NOT Be Merged

To prevent technical debt, breaking changes, and architectural degradation, the following boundaries **must remain strictly separated**:

1. **DO NOT merge RAG models into Foundation Orchestration persistence models**:
   - rag_documents and rag_chunks must remain in rag/storage/models.py.
   - Do NOT attempt to convert ChunkModel into a generic DataReference or ArtifactReference. Chunks require pgvector column types, specific vector indices, and chunk metadata that orchestrator entities do not possess.
2. **DO NOT rewrite RAG to depend on FoundationCore or AppContext**:
   - RAG must remain an independent engine that can be run standalone from CLI or imported as a clean library.
3. **DO NOT replace Docling in RAG with Foundation's DoclingDocumentParser**:
   - RAG's chunker depends specifically on the markdown structural hierarchies emitted by DoclingDocumentIngester. Foundation's parser outputs flat ExtractedTable / ExtractedFigure structures meant for JSON payload inspection.
4. **DO NOT inject LLM generation directly into the retrieval or reranker layers**:
   - Retrieval and reranking must return raw candidates and scores. LLM answer generation is the responsibility of downstream agents/workflows.

5. **DO NOT bypass the connector seam**:
    - Never allow orchestration/planning/ or apps/api/routers/ to perform ad-hoc SQL vector searches against rag_chunks. All retrieval must pass through RetrievalConnector.

---

16. Integration Architecture Map

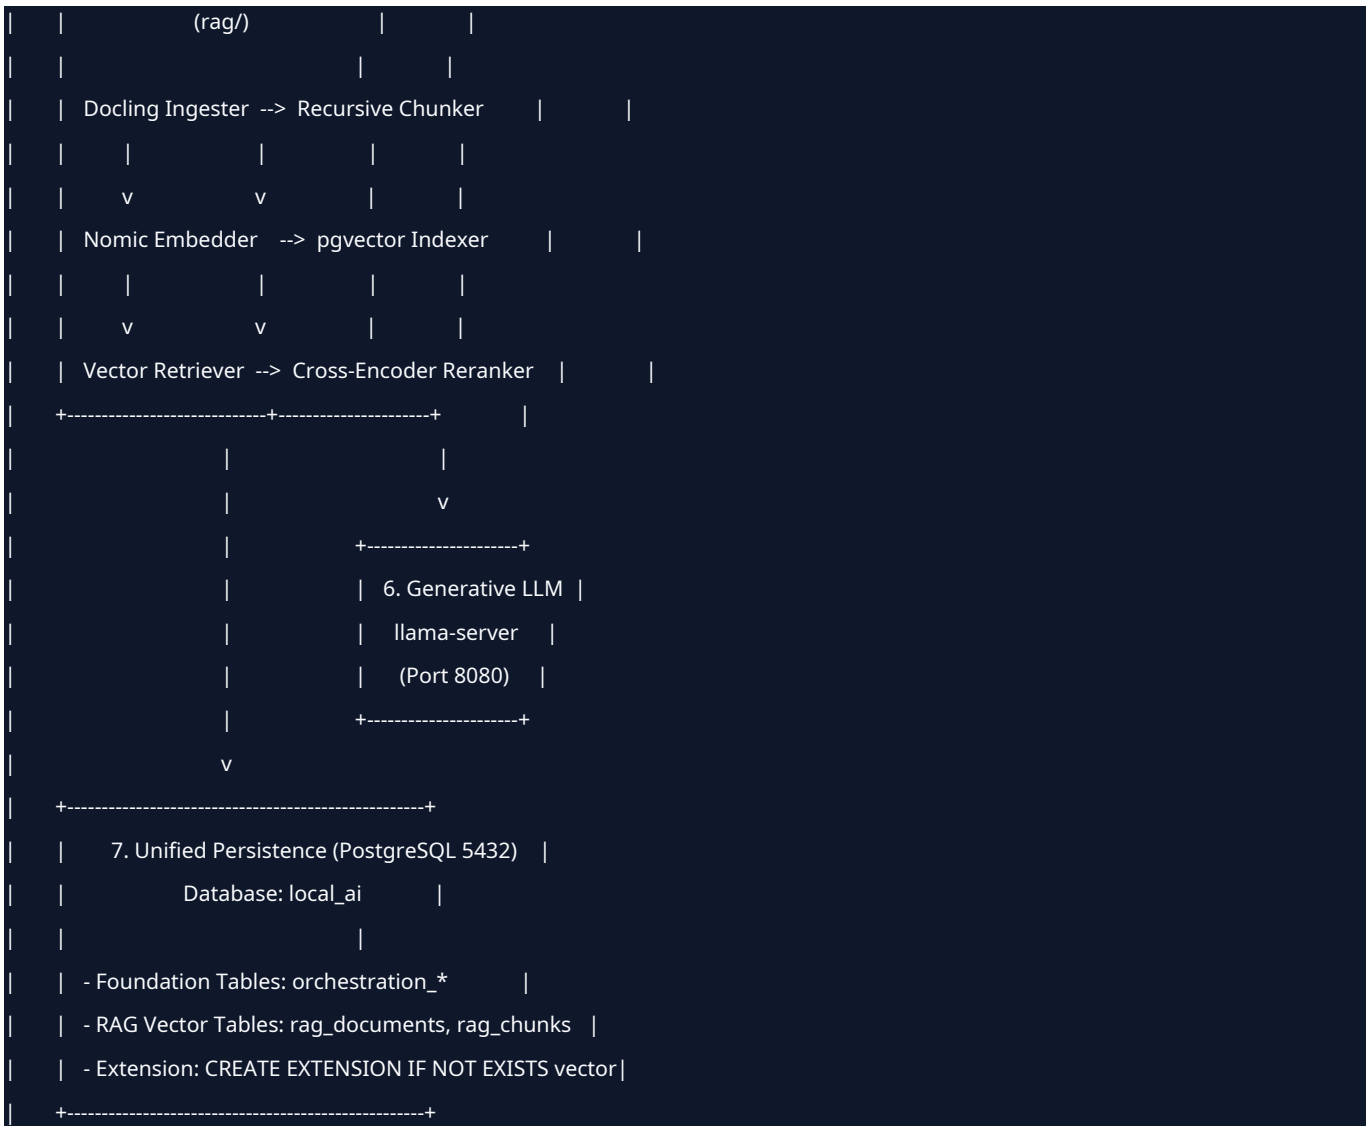

---

17. Integration Readiness Assessment

Overall System Readiness: YELLOW (Architecture Ready, Environment Blockers Present)

#	Subsystem / Area	Status	Key Diagnostic & Blocker Summary
1	**RAG Core Pipeline**	**GREEN**	Fully implemented, 135 passing unit tests, zero network calls, robust error handling.
2	**RAG CLI Test Harness**	**GREEN**	Fully operational with rich console output; document ID display bug already resolved.
3	**RAG Offline**	**GREEN**	Full zero-internet guard verified; all 5 HF cache models present

#	Subsystem / Area	Status	Key Diagnostic & Blocker Summary
	Hardening		locally on disk.
4	**Inference Architecture**	**YELLOW**	Design is sound and non-conflicting, but `llama-server` is not built and GGUF models are missing.
5	**PostgreSQL Persistence**	**YELLOW**	Tables are completely non-colliding, but `.env` and `settings.toml` point to divergent DB names (`ragdb` vs `local_ai`).
6	**Database Migrations**	**YELLOW**	Foundation has Alembic `001`; RAG uses declarative `create_all()`. Needs unified migration.
7	**FastAPI Delivery Layer**	**RED**	Missing `fastapi` and `uvicorn` in `.venv`; cannot launch HTTP server.
8	**Workspace Sandbox**	**RED**	Missing `docker` package and Docker daemon command; code repair workflow will fail.
9	**Orchestration Runner**	**GREEN**	`InProcessTaskRunner` operates cleanly without external Temporal cluster.
10	**Configuration Unification**	**YELLOW**	Settings split across `settings.toml`, `llama_models.ini`, and `.env`.
11	**Dependency Specifications**	**RED**	`pyproject.toml` lacks RAG dependencies; `.venv` lacks Foundation dependencies.
12	**Automated Test Harness**	**YELLOW**	`pytest` is not installed in `.venv`; `pyproject.toml` ignores `rag/tests/`.
13	**RAG Foundation Seam**	**GREEN**	Connector and Capability architectures are cleanly defined; zero invasiveness required.

---

18. Phased Recommended Merge Order (Roadmap for Phase 2)

```
Phase 2.1: Environment & Dependency Harmonization
   ├─── Update pyproject.toml (add pgvector, torch, transformers, sentence-transformers, rich, python-dotenv)
   ├─── Install missing runtime packages (fastapi, uvicorn, pytest) into .venv
   └─── Update pyproject.toml testpaths to include both ["tests", "rag/tests"]

Phase 2.2: Database & Configuration Consolidation
   ├─── Align .env and settings.toml to point to single database 'local_ai'
   ├─── Create Alembic migration 002_rag_vector_schema.py to formalize rag_* tables
   └─── Run alembic upgrade head to ensure pgvector extension and all tables exist in local_ai

Phase 2.3: Foundation Offline Hardening
   ├─── Apply rag/offline.py environment guard to orchestration/capabilities/builtin/document/docling_parser.py
   └─── Add enable_remote_services=False to Foundation Docling converters

Phase 2.4: Architectural Seam Implementation
   ├─── Create connectors/retrieval.py defining RetrievalConnector interface and LocalRAGRetrievalConnector
   ├─── Create orchestration/capabilities/builtin/retrieval/ exposing 'retrieval.search' and 'rag.ingest'
   └─── Register RetrievalConnector into apps/context.py AppContext

Phase 2.5: API Surface Extension
   ├─── Add /api/v1/direct/retrieve endpoint to apps/api/routers/direct.py
   └─── Add /api/v1/rag/ingest endpoint for document upload + embedding indexing

Phase 2.6: Verification & Full-System Smoke Testing
   ├─── Run full test suite: pytest tests/ rag/tests/
   └─── Execute end-to-end integration test: Upload PDF -> Ingest/Embed -> Vector Retrieval -> Rerank
```

Audit completed with zero modifications to source files.